

LinearMDP Likelihood Testing

MDP parameters = {'n':2, 'b':1, 'determinism':1.0, 'discount':0.99, 'seed': 0}

Same paths and true R

Matlab Reward Function Output

```
True R
0.000528330112439 0.000528330112439 0.000528330112439 0.000528330112439 0.000528330112439
0.000000002399272 0.000000002399272 0.000000002399272 0.000000002399272 0.000000002399272
4.510928037168985 4.510928037168985 4.510928037168985 4.510928037168985 4.510928037168985
4.533927110477543 4.533927110477543 4.533927110477543 4.533927110477543 4.533927110477543

True Likelihood
1.043640655503507e+04 = 10436.64

Optimal Policy
2
2
3
1

Found R
-2.282269456199177 -2.282269456199177 -2.282269456199177 -2.282269456199177 -2.282269456199177
-1.630248097698442 -1.630248097698442 -1.630248097698442 -1.630248097698442 -1.630248097698442
2.853838697068992 2.853838697068992 2.853838697068992 2.853838697068992 2.853838697068992
2.874558020258835 2.874558020258835 2.874558020258835 2.874558020258835 2.874558020258835

Found Likelihood
1.043568101290607e+04 = 10435.68

Optimal Policy
2
2 = [1,1,2,0] in python
3 cos index diff
1
```

Numpy Reward Function Output

```
True R is
[[0.0005 0.0005 0.0005 0.0005 0.0005]
 [0.      0.      0.      0.      0.    ]
 [4.5109 4.5109 4.5109 4.5109 4.5109]
 [4.5339 4.5339 4.5339 4.5339 4.5339]]
with negated likelihood of 10436.406555
and optimal policy [1 1 2 0]

Found R is
[[-2.73117114 -2.73117114 -2.73117114 -2.73117114 -2.73117114]
 [-2.05036447 -2.05036447 -2.05036447 -2.05036447 -2.05036447]
 [ 2.43509587  2.43509587  2.43509587  2.43509587  2.43509587]
 [ 2.45579707  2.45579707  2.45579707  2.45579707  2.45579707]]
with negated likelihood of 10435.6824895
and optimal policy [1 1 2 0]
```

PyTorch Reward Function Output

```
True R is
tensor([[5.000e-04, 5.000e-04, 5.000e-04, 5.000e-04, 5.000e-04],
        [0.000e+00, 0.000e+00, 0.000e+00, 0.000e+00, 0.000e+00],
        [4.511e+00, 4.511e+00, 4.511e+00, 4.511e+00, 4.511e+00],
        [4.534e+00, 4.534e+00, 4.534e+00, 4.534e+00, 4.534e+00]])
with negated likelihood of 10436.40655548196
and optimal policy [1 1 2 0]

Found R is
tensor([[ -3.024,  -3.024,  -3.024,  -3.024,  -3.024],
        [-2.375, -2.375, -2.375, -2.375, -2.375],
        [ 2.118,  2.118,  2.118,  2.118,  2.118],
        [ 2.139,  2.139,  2.139,  2.139,  2.139]], dtype=torch.float64)
with negated likelihood of 10435.681479438554
and optimal policy [1 1 2 0]
```

PyTorch with custom sampled paths using linear NN w/out backward function

```
EstR: tensor([[ -1.620,  -1.458,   3.217,   3.242]], grad_fn=<AddmmBackward>) | loss: 10436.089196453906 | EVD: 0.00283340380700825
```

Matlab MDP Solution Output

```
P
0.002880150863316    0.988350220895028    0.003009326515067    0.002880150863316    0.002880150863316
0.002877354167903    0.988614094053452    0.002877354167903    0.002877354167903    0.002753843442797
0.246970950890849    0.246970950890849    0.258367449408286    0.000719697919203    0.246970950890849
0.252601744935846    0.252601744935846    0.252601744935846    0.000735195551007    0.241459569641488

Q
1.0e+02 *

5.791843647590228    5.850225591926778    5.792282383877290    5.791843647590228    5.791843647590228
5.792277100600159    5.850671429371738    5.792277100600159    5.792277100600159    5.791838364313096
5.895329588997344    5.895329588997344    5.895780709719435    5.836947644660794    5.895329588997344
5.896010700452520    5.896010700452520    5.896010700452520    5.837616371680941    5.895559579730429

V
1.0e+02 *

5.850342773621333
5.850785941588063
5.909314434567343
5.909770112064405
```

Python MDP Solution Output

```
P
[[ 0.00288015  0.98835014  0.00300941  0.00288015  0.00288015]
 [ 0.00287743  0.98861386  0.00287743  0.00287743  0.00275384]
 [ 0.24697086  0.24697086  0.25836772  0.0007197  0.24697086]
 [ 0.25260182  0.25260182  0.25260182  0.00073522  0.24145931]]

Q
[[ 579.18159787  585.01979289  579.22550107  579.18159787  579.18159787]
 [ 579.22500107  585.06440635  579.22500107  579.22500107  579.18109787]
 [ 589.53019291  589.53019291  589.57530637  583.69199789  589.53019291]
 [ 589.59830614  589.59830614  589.59830614  583.75890085  589.55319267]]

V
[[ 585.03151114]
 [ 585.07585781]
 [ 590.92867782]
 [ 590.97424698]]
```

Scaling R and recalculating with same paths in Python:

```
True R is
[[ 0.82299711  0.82299711  0.82299711  0.82299711  0.82299711]
 [ 6.84492019  6.84492019  6.84492019  6.84492019  6.84492019]
 [ 1.74250829  1.74250829  1.74250829  1.74250829  1.74250829]
 [ 0.77701112  0.77701112  0.77701112  0.77701112  0.77701112]]
with negated likelihood of 7926.74335351
and optimal policy [2 0 3 3]

... Doubling R & recalcuating ...

Double R is
[[ 1.64599423  1.64599423  1.64599423  1.64599423  1.64599423]
 [13.68984039 13.68984039 13.68984039 13.68984039 13.68984039]
 [ 3.48501658  3.48501658  3.48501658  3.48501658  3.48501658]
 [ 1.55402224  1.55402224  1.55402224  1.55402224  1.55402224]]
with negated likelihood of 7926.74335351
and optimal policy [2 0 3 3]

... Quadrupling R & recalculating ...

Quadrupled R is
[[ 3.29198845  3.29198845  3.29198845  3.29198845  3.29198845]
 [27.37968078 27.37968078 27.37968078 27.37968078 27.37968078]
 [ 6.97003317  6.97003317  6.97003317  6.97003317  6.97003317]
 [ 3.10804448  3.10804448  3.10804448  3.10804448  3.10804448]]
with negated likelihood of 7926.74335351
and optimal policy [2 0 3 3]
```